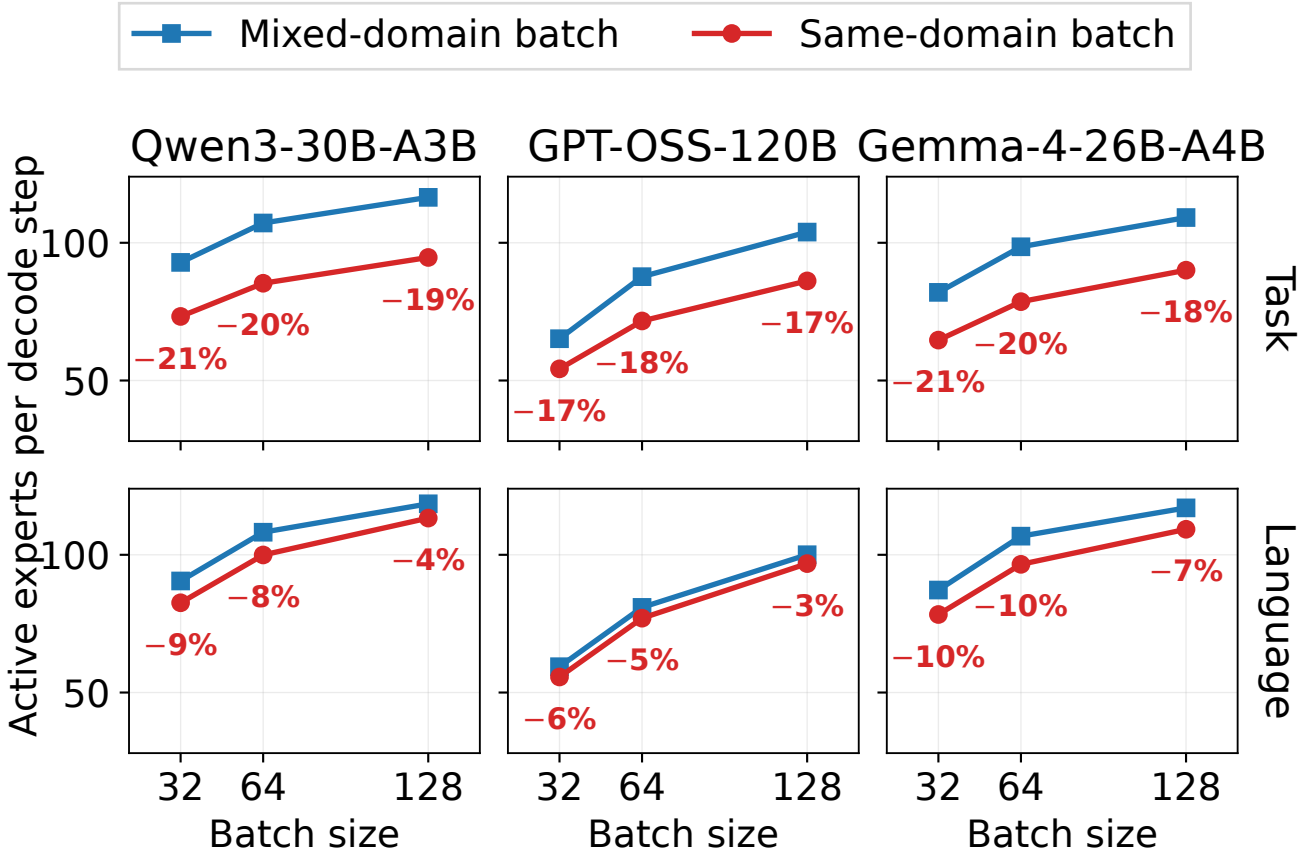
Mixed-domain batch
Same-domain batch
Qwen3-30B-A3B
GPT-OSS-120B
Gemma-4-26B-A4B
Task
Language
Active experts per decode step
100
50
−21%
−20%
−19%
−17%
−18%
−17%
−21%
−20%
−18%
−9%
−8%
−4%
−6%
−5%
−3%
−10%
−10%
−7%
32
64
128
Batch size