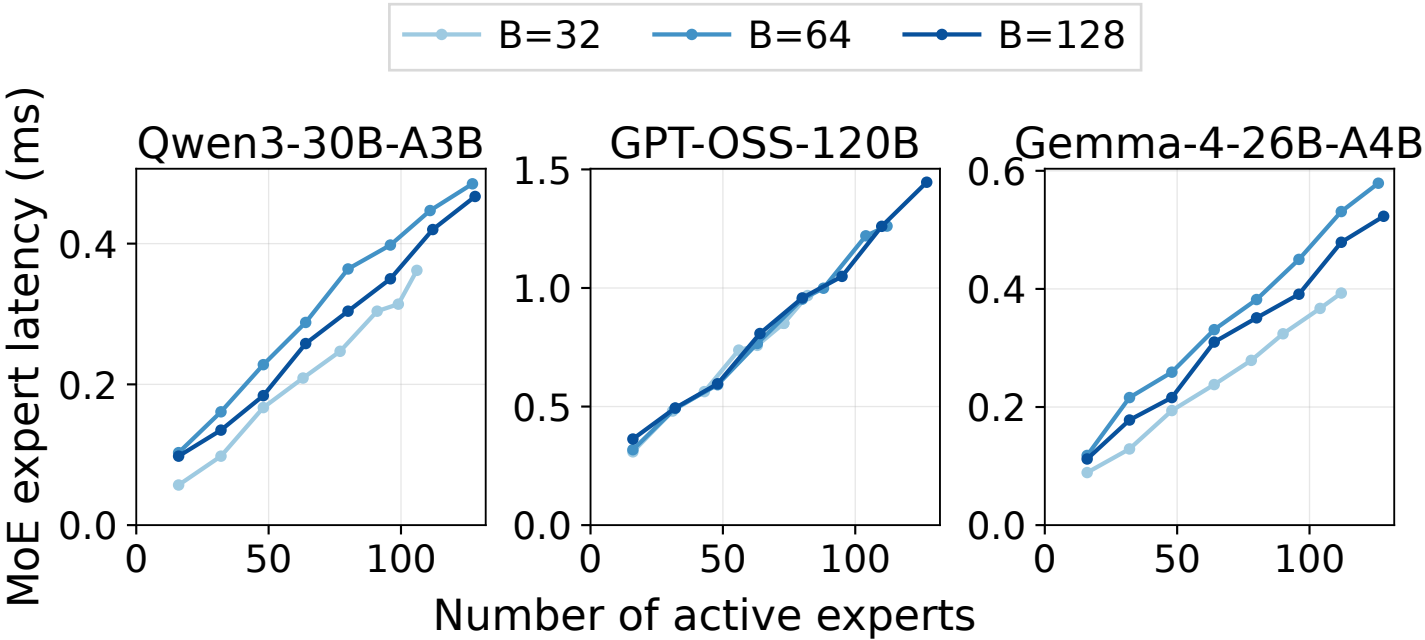
B=32
B=64
B=128
Qwen3-30B-A3B
GPT-OSS-120B
Gemma-4-26B-A4B
MoE expert latency (ms)
Number of active experts
0.0
0.2
0.4
0.5
1.0
1.5
0.6
0
50
100